


General Agent Problem Solving

Architecting the Future of AI-Native Companies



| The Paradigm Shift: **Labor vs Tools**

- ✓ **SaaS Era (Software Tools):** Focus on human-operated digital toolkits.
- ✓ **Agent Factory Era (Labor):** Focus on autonomous AI Employees.
- ✓ **Value Metric Shift:** Moving from "Per-Seat" subscriptions to "Per-Outcome" results.
- ✓ **Visible vs. Industrialized:** Moving from manual tasking to automated high-fidelity artifacts.



| Defining **GAPS** Framework



Intent

High-fidelity blueprinting of goals, constraints, budgets, and specific permissions.



Production

The engine that transforms raw intent into results via industrialized AI station workflows.



Result

Verified actions and artifacts delivered on-demand and improved via continuous feedback loops.

| The 10-80-10 Rhythm

10%

HUMAN INTENT

Clear planning, budget, and limits.

80%

AI EXECUTION

Industrialized production of
work.

10%

VERIFICATION

Human review and final
approval.

Edge layer

Serves the
individual



Personal identic agent

OpenClaw · your chief of staff



Knows your
context




Speaks on
your behalf



Acts within
your authority

| Action & Precision: Principles 1-2



1. Bash is the Key

The agent does not just describe a solution; it acts. It executes commands to manipulate the environment.

2. Code as Universal Interface

Precision comes from structured formats—schemas, tables, and code blocks—not ambiguous prose.


| Reliability & Structure: Principles 3-4

3. Verification as Core Step

Every meaningful output is checked before shipping.
"Looks right" is considered a failure mode.

4. Small, Reversible Decomposition

Work proceeds in atomic steps. Every step can be undone or reversed, minimizing catastrophic failures.



| Memory & Constraints: Principles 5-6



5. Persisting State in Files

Conversations are volatile; filesystems are durable. All critical context lives in a structured directory.


6. Constraints and Safety

Explicit permissions and scope. Autonomy is earned per task type, never granted by default.

Principle 7: Observability

No black boxes, no surprises. The principal must be able to trace every action the agent takes.

- ✓ **Full Audit Trail:** Logs for every CLI command.
- ✓ **Real-time Monitoring:** Visibility into the "thought" process.
- ✓ **Transparency:** Easy to debug and verify intent alignment.



Observability platform

| The Three Invariants

- ✓ **The Human Principal:** Every legitimate action originates with a human who owns the outcome and budget.
- ✓ **The Delegate:** A personal agent that holds context, judgment, and authority to broker downstream work.
- ✓ **The Management Layer:** The OS of the AI-native company—hires workers, assigns tasks, and enforces policies.




| Problem Solving vs Manufacturing



Category	Problem Solving	Manufacturing
Objective	Immediate Answer	Permanent Worker
Scope	Fix a bug / Analyze	Recurring Workflows
Outcome	Session ends	Scalable Labor

| The **AI-Native** Architecture

- ✓ **Right Engine:** Reliable engines for critical work; cost-efficient for routines.
- ✓ **System of Record:** Official company memory for reading and writing.
- ✓ **Lifecycle Management:** Automated hiring, tasking, and retiring of AI workers.
- ✓ **Traceability:** Every outcome is verified and owned by a human principal.



Thank You